



PURDUE CS47100
SEPT 9, 2019
PROF. JENNIFER NEVILLE

INTRODUCTION TO AI

ADVERSARIAL GAMES




BEHAVIOR FROM COMPUTATION




TYPES OF GAMES

- ▶ Many different kinds of games
- ▶ Axes:
 - ▶ Deterministic vs. stochastic
 - ▶ One, two, or more players
 - ▶ Zero sum
 - ▶ Perfect information (can you see the state)
- ▶ Algorithms need to calculate a strategy (**policy**) which recommends a move from each state




DETERMINISTIC GAMES

- ▶ Problem formulation:
 - ▶ States: S (start at s_0)
 - ▶ Players: $P=\{1\dots N\}$ (usually take turns)
 - ▶ Actions: A (may depend on player / state)
 - ▶ Transition Function: $S \times A \rightarrow S$
 - ▶ Terminal Test: $S \rightarrow \{t, f\}$
 - ▶ Terminal Utilities: $S \times P \rightarrow R$
- ▶ Solution for a player is a **policy**: $S \rightarrow A$




ZERO-SUM GAMES



▶ Zero-Sum Games


- ▶ Agents have opposite *utilities* (values on outcomes)
- ▶ Can then think of outcome as a single value that one maximizes and the other minimizes
- ▶ Adversarial, pure competition




▶ General Games

- ▶ Agents have independent *utilities* (values on outcomes)
- ▶ Cooperation, indifference, competition, and more are all possible
- ▶ More later on non-zero-sum games

ADVERSARIAL SEARCH

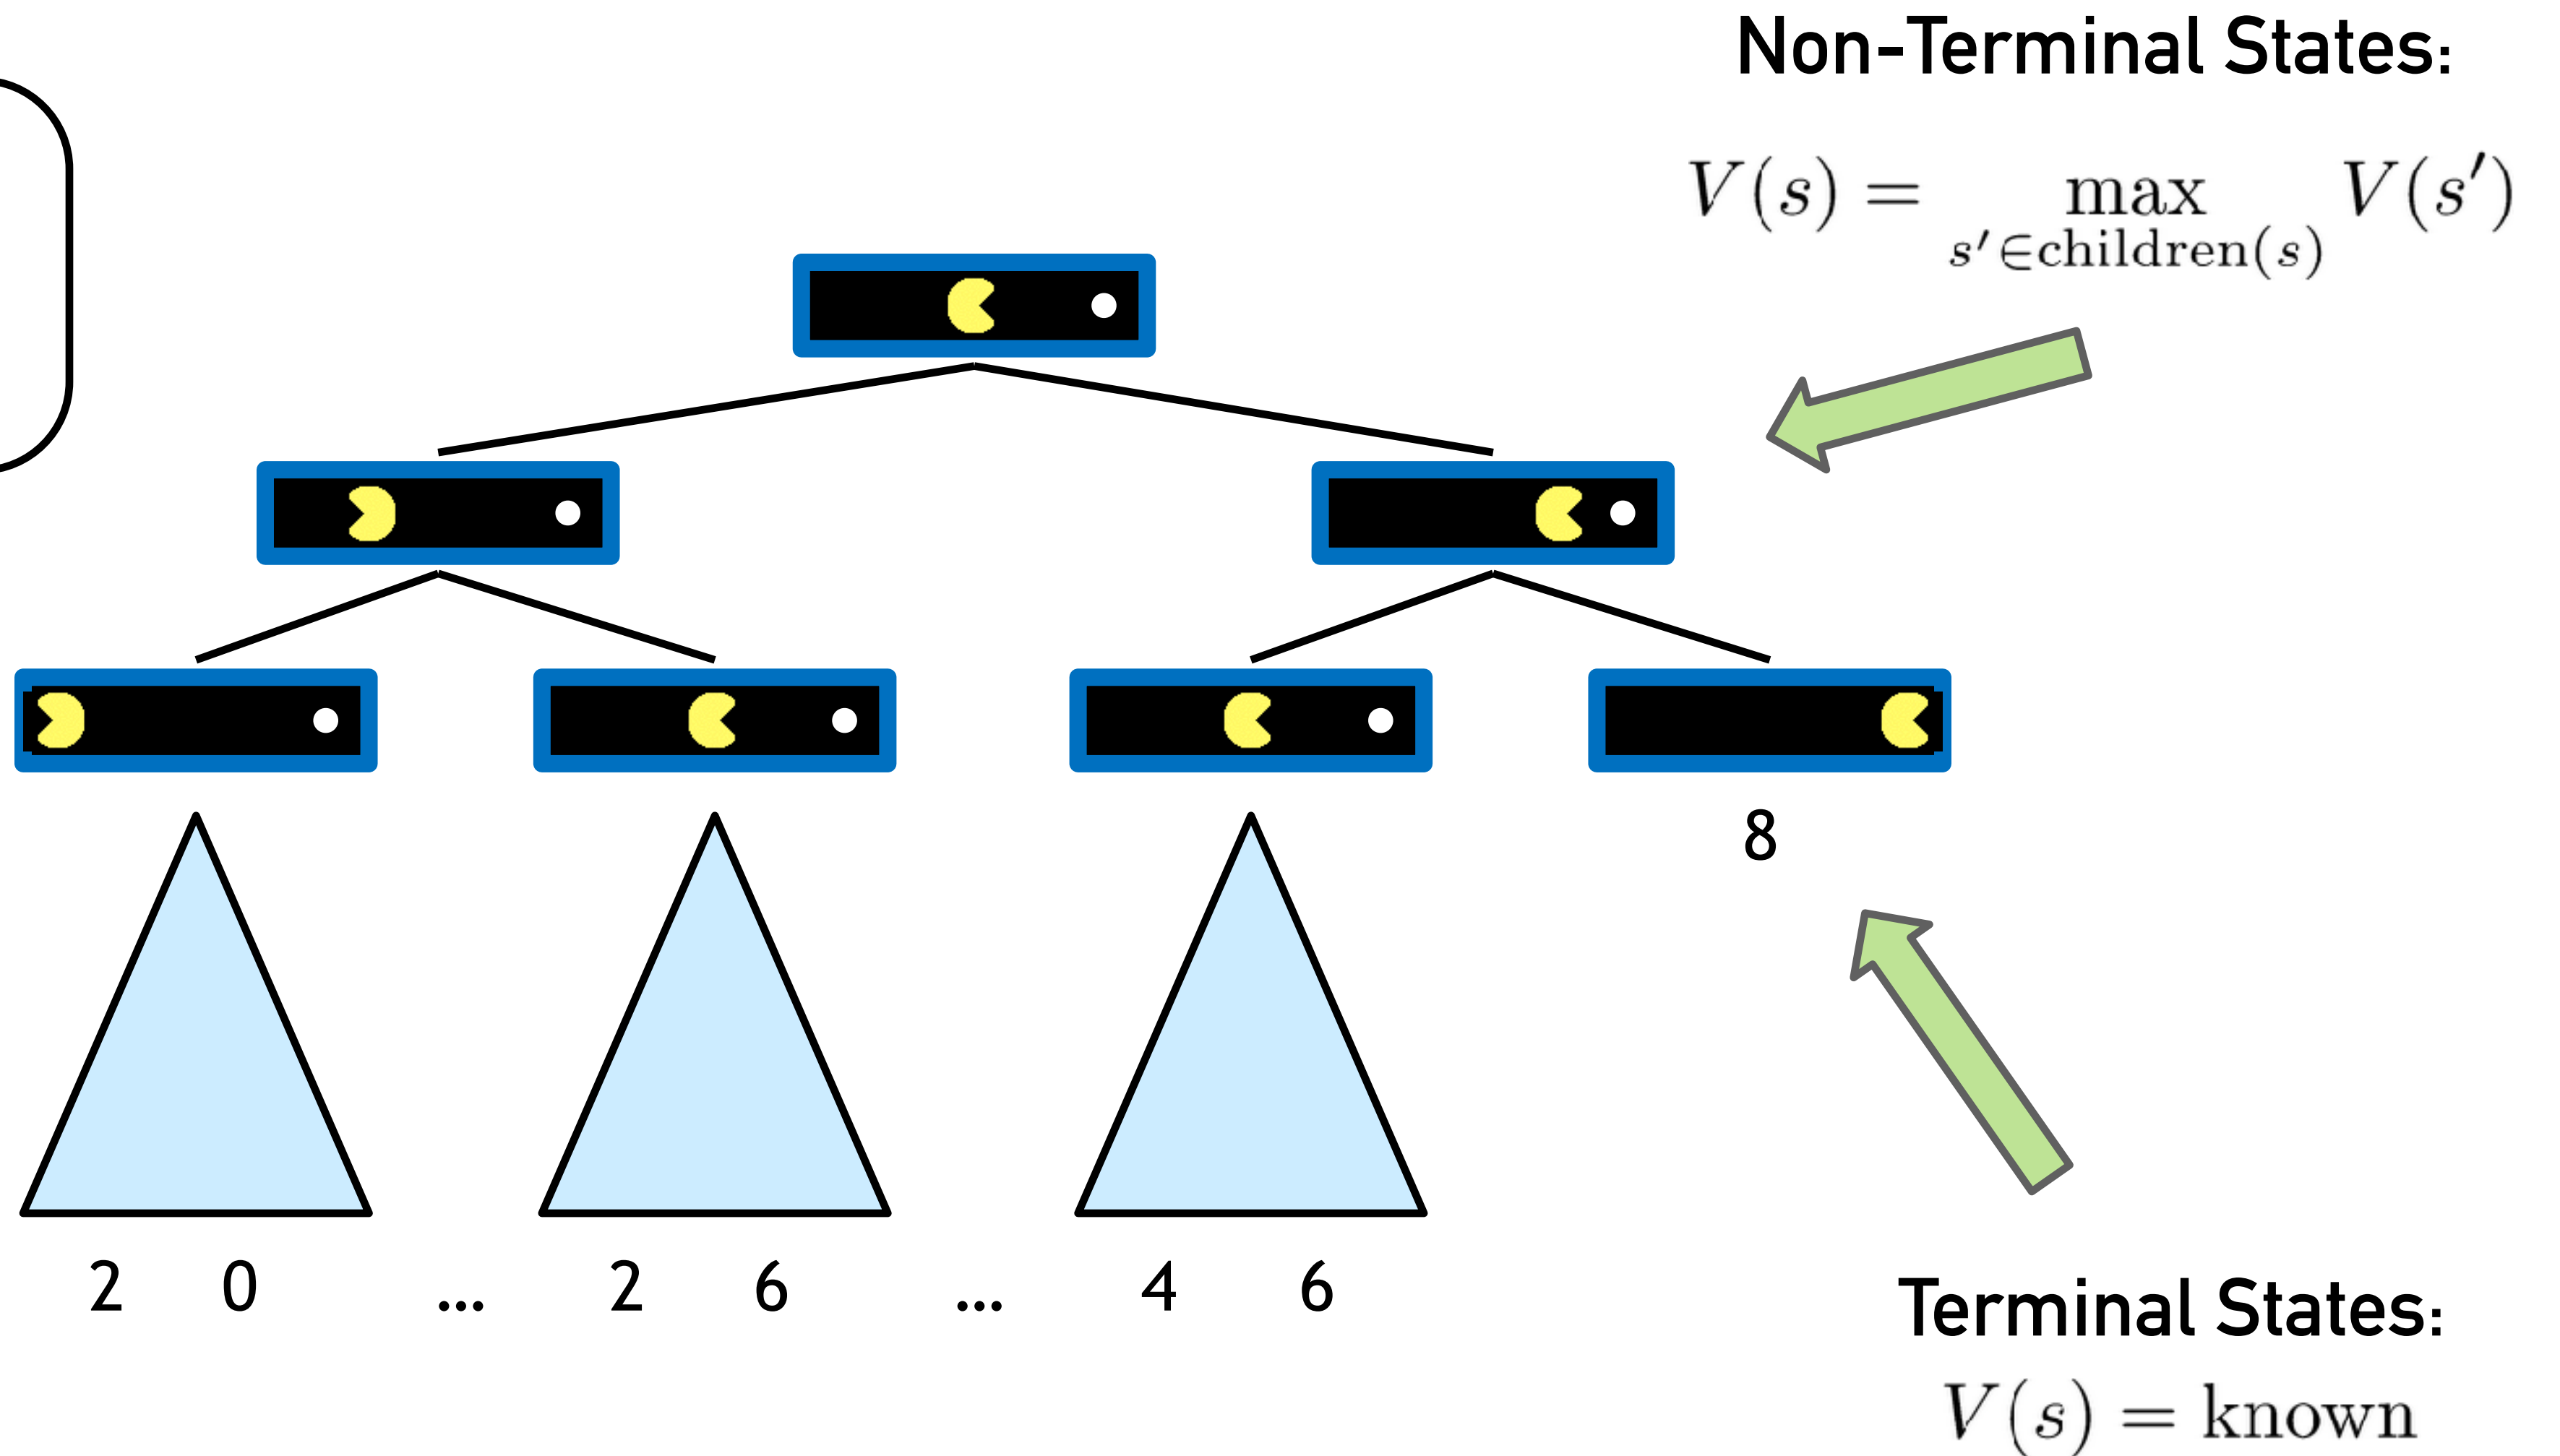


SINGLE-AGENT TREES




VALUE OF A STATE

Value of a state:
The best achievable
outcome (utility) from
that state



ADVERSARIAL GAME TREES




MINIMAX VALUES

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:

$$V(s) = \text{known}$$

TIC-TAC-TOE GAME TREE



MAX (X)



MIN (O)





MAX (X)



MIN (O)

TERMINAL

Utility

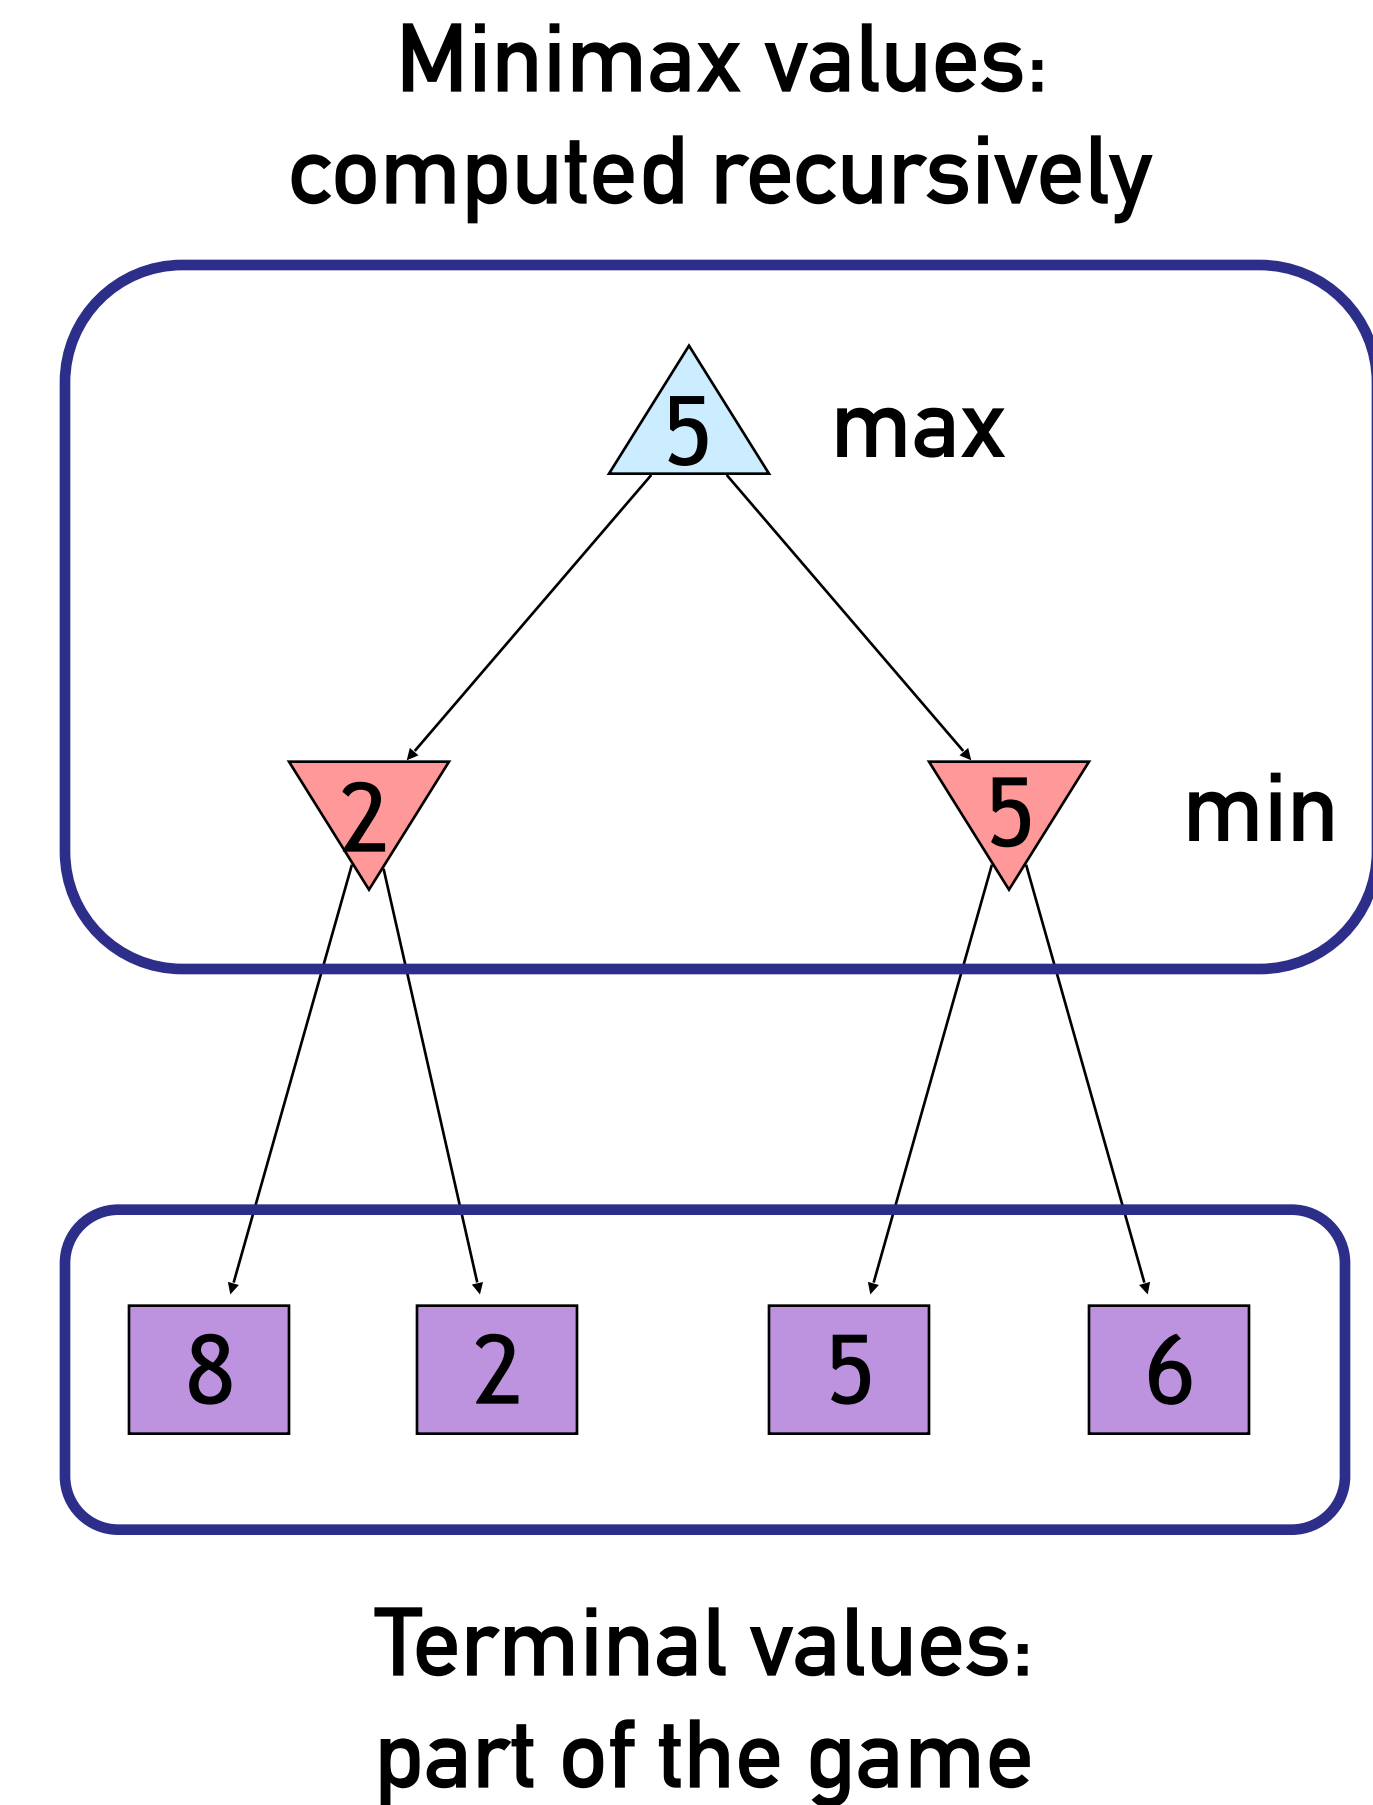


ADVERSARIAL SEARCH (MINIMAX)

- ▶ Deterministic, zero-sum games:
 - ▶ Tic-tac-toe, chess, checkers
 - ▶ One player maximizes result
 - ▶ The other minimizes result

- ▶ Minimax search:


- ▶ A state-space search tree
- ▶ Players alternate turns
- ▶ Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary



MINIMAX IMPLEMENTATION

```
def max-value(state):  
    initialize v =  $-\infty$   
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$



```
def min-value(state):  
    initialize v =  $+\infty$   
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```


$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

MINIMAX IMPLEMENTATION (DISPATCH)

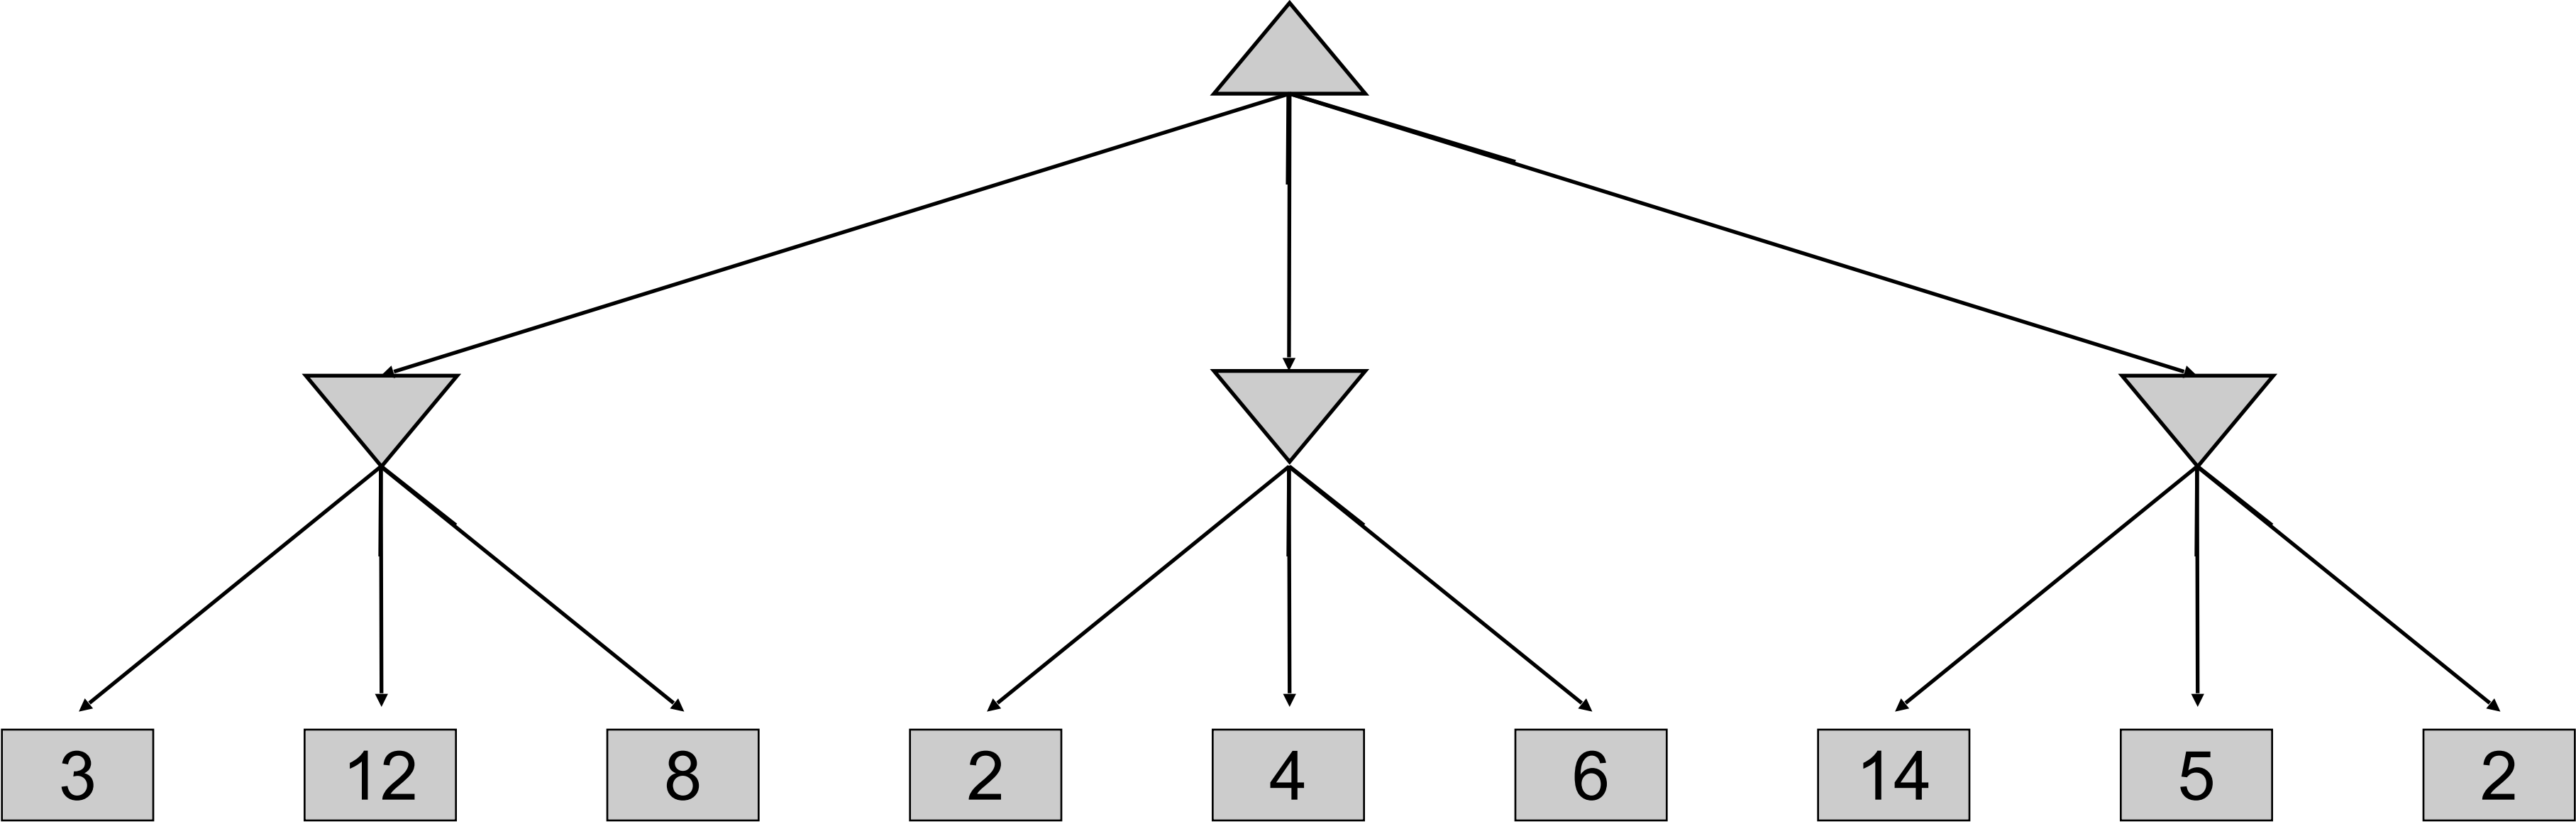
```
def value(state):  
    if the state is a terminal state: return the  
        state's utility  
    if the next agent is MAX: return max-value(state)  
    if the next agent is MIN: return min-value(state)
```

```
def max-value(state):  
    initialize v =  $-\infty$   
    for each successor of state:  
        v = max(v, value(successor))  
    return v
```




```
def min-value(state):  
    initialize v =  $+\infty$   
    for each successor of state:  
        v = min(v, value(successor))  
    return v
```



MINIMAX EXAMPLE




MINIMAX PROPERTIES



Optimal against a perfect player. Otherwise?

MINIMAX EFFICIENCY

- ▶ Efficient of minimax search
 - ▶ Just like (exhaustive) DFS
 - ▶ Time: $O(b^m)$
 - ▶ Space: $O(bm)$
- ▶ Example: For chess, $b \approx 35$, $m \approx 100$
 - ▶ Exact solution is completely infeasible
 - ▶ But, do we need to explore the whole tree?




GAME TREE SIZES

- ▶ Tic-tac-toe: 10^5
- ▶ Checkers: 10^{31}
- ▶ Chess: 10^{123}
- ▶ Backgammon: 10^{144}
- ▶ Go: 10^{360}

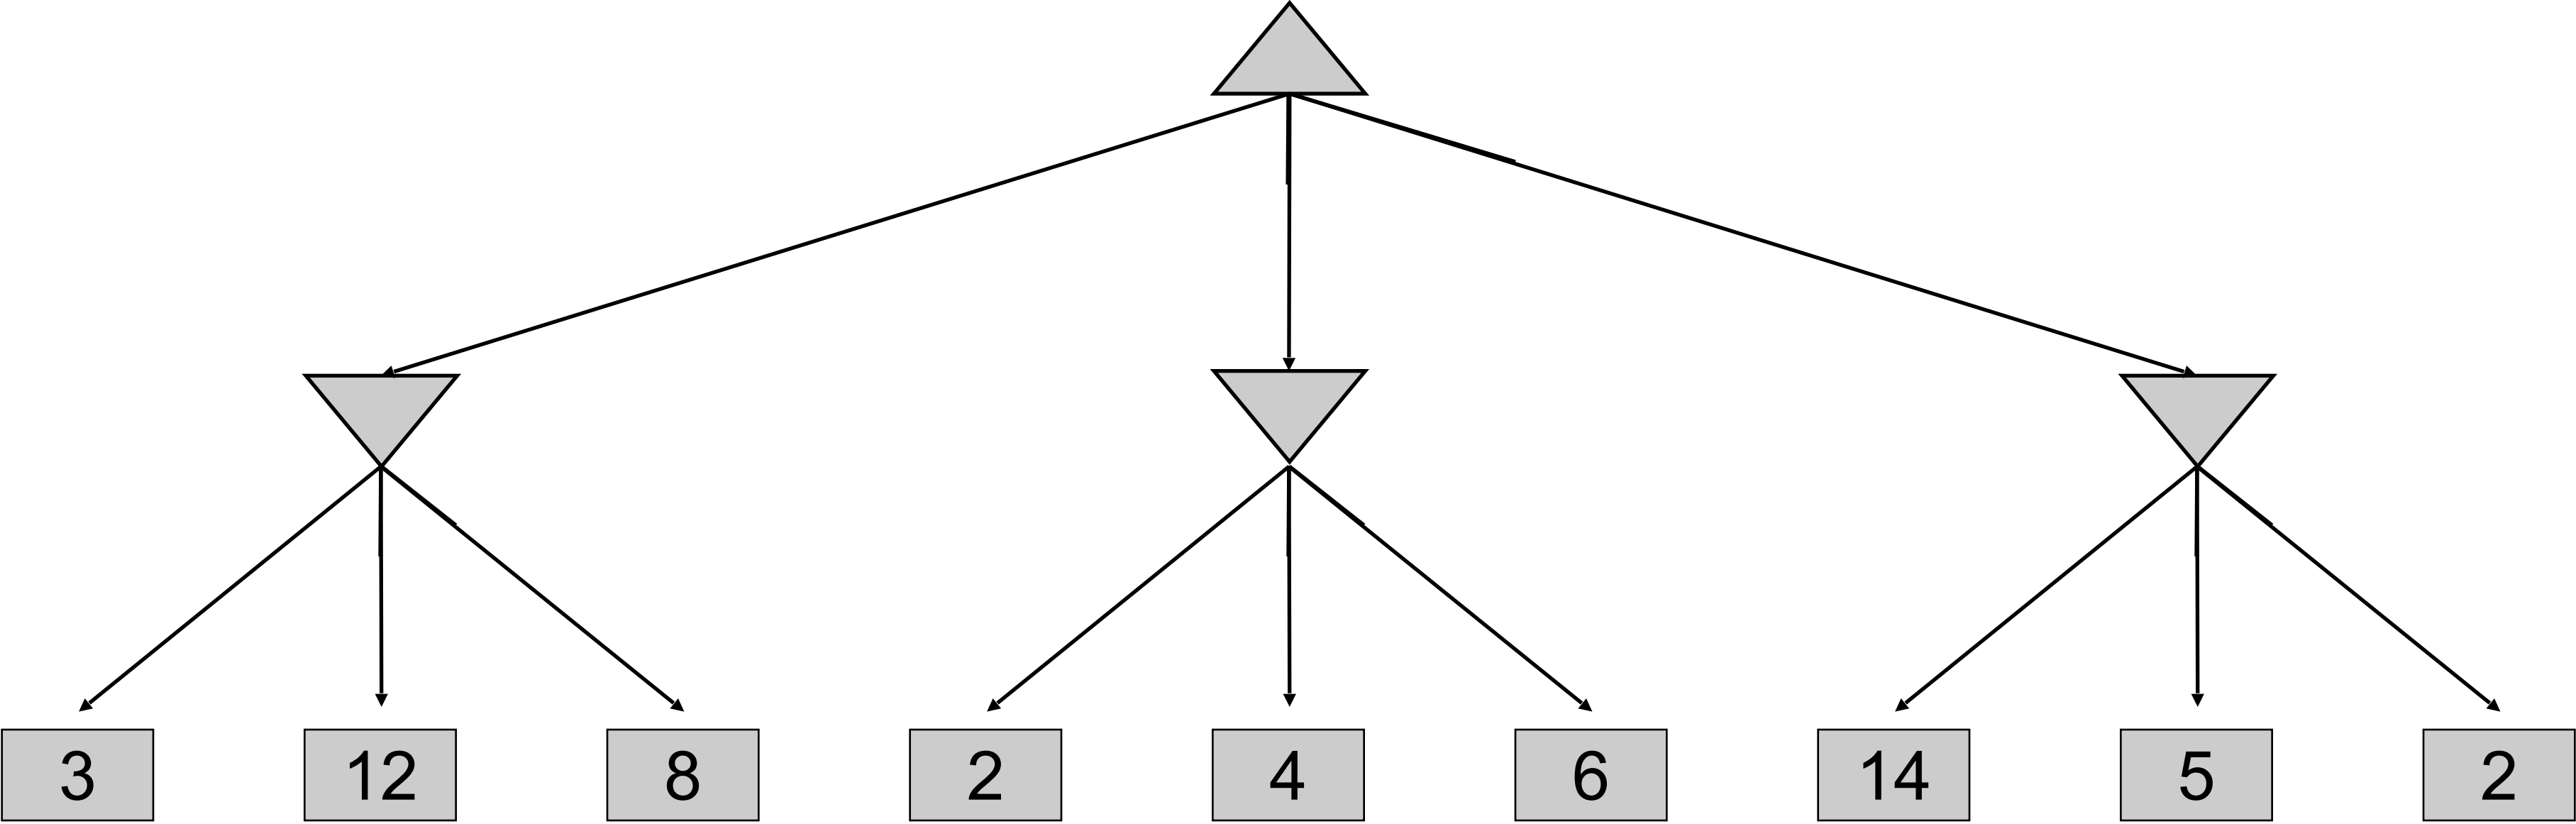
Assume that a computer can evaluate 1 million board configurations per second.

Then it would take **0.1 seconds to search** the entire tic-tac-toe game tree but it would still take **10^{18} years** to search the full checkers tree.


GAME TREE PRUNING



MINIMAX EXAMPLE




MINIMAX PRUNING



ALPHA-BETA PRUNING

- ▶ General configuration (MIN version)
 - ▶ When computing the MIN-VALUE at some node n
 - ▶ We loop over n 's children
 - ▶ n 's estimate of the children's min is dropping
 - ▶ Who cares about n 's value? MAX
 - ▶ Let a be the best value that MAX can get at any choice point along the current path from the root
 - ▶ If n becomes worse than a , MAX will avoid it, so we can stop considering n 's other children (it's already bad enough that it won't be played)
- ▶ MAX version is symmetric



Alpha-Beta Implementation


α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

ALPHA-BETA PRUNING PROPERTIES

- ▶ This pruning has **no effect** on minimax value computed for the root
- ▶ Values of intermediate nodes might be wrong
 - ▶ Important: children of the root may have the wrong value
 - ▶ So the most naïve version won't let you do action selection
- ▶ Good child ordering improves effectiveness of pruning
- ▶ With "perfect ordering":
 - ▶ Time complexity drops to $O(b^{m/2})$
 - ▶ Doubles solvable depth
 - ▶ Full search of, e.g. chess, is still hopeless...
- ▶ This is a simple example of **metareasoning** (computing about what to compute)




RESOURCE LIMITS





RESOURCE LIMITS

- ▶ **Problem:** In realistic games, cannot search to leaves
- ▶ **Solution:** Depth-limited search
 - ▶ Instead, search only to a limited depth in the tree
 - ▶ Replace terminal utilities with an evaluation function for non-terminal positions
- ▶ **Example:**
 - ▶ Suppose we have 100 seconds, can explore 10K nodes / sec
 - ▶ So can check 1M nodes per move
 - ▶ α - β reaches about depth 8 - decent chess program
- ▶ Guarantee of optimal play is gone
- ▶ More plies makes a BIG difference
- ▶ Use iterative deepening for an anytime algorithm




DEPTH MATTERS

- ▶ Evaluation functions are always imperfect
- ▶ The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- ▶ An important example of the tradeoff between complexity of features and complexity of computation




EVALUATION FUNCTIONS



EVALUATION FUNCTIONS

- ▶ Evaluation functions score non-terminals in depth-limited search



- ▶ Ideal function: returns the actual minimax value of the position
- ▶ In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.