


PURDUE CS47100
SEPT 11, 2019
PROF. JENNIFER NEVILLE

INTRODUCTION TO AI

RECAP: ADVERSARIAL SEARCH

- ▶ Minimax search is a way of finding an optimal move in a zero-sum two player game
- ▶ Alpha-beta pruning is a way of finding the optimal minimax solution while avoiding searching subtrees of moves which won't be selected
 - ▶ Some branches will never be played by rational players since they include sub-optimal decisions (for either player)
 - ▶ Pruning produces results that are exactly equivalent to complete (unpruned) search
 - ▶ Node *ordering* can improve effectiveness; Perfect ordering gives time complexity $O(b^{m/2})$, thus, can search twice as far as ordinary minimax in equal time




RESOURCE LIMITS




RESOURCE LIMITS

- ▶ **Problem:** In realistic games, cannot search to leaves
- ▶ Example:
 - ▶ Suppose we have 100 seconds, can explore 10K nodes / sec
 - ▶ So can check 1M nodes per move
 - ▶ α - β reaches about depth 8 - decent chess program
- ▶ Guarantee of optimal play is gone; More plies makes a BIG difference
- ▶ **Solution:** Depth-limited search
 - ▶ Instead, search only to a limited depth in the tree
 - ▶ Replace terminal utilities with an **evaluation function** for non-terminal positions




WHAT TO DO WHEN SEARCH IS INTRACTABLE



- ▶ Stop the search before you reach terminal states (using depth cutoff)
- ▶ Evaluate nodes using an evaluation function – What properties should the evaluation function have?

EVALUATION FUNCTIONS



EVALUATION FUNCTIONS

- ▶ Desirable properties
 - ▶ Order terminal states in same way as true utility function
 - ▶ Strongly correlated with the actual minimax value of the states
 - ▶ Efficient to calculate
- ▶ Typically use **features** – simple characteristics of the game state that are correlated with the *probability of winning*
- ▶ The evaluation function combines feature values to produce a score:



$$Eval(x) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s) = \sum_{i=1}^n w_i f_i(s)$$

EXAMPLE FEATURES

- ▶ What would be some useful features for chess?
 - ▶ Relative number of Bishops; Knights; Rooks; Pawns
 - ▶ Total number of pieces
 - ▶ Has queen?
 - ▶ Castled?
 - ▶ In check?
 - ▶ Distance of furthest pawn from start
 - ▶ Relative freedom (relative total number of possible moves)
 - ▶ etc.

DEPTH MATTERS

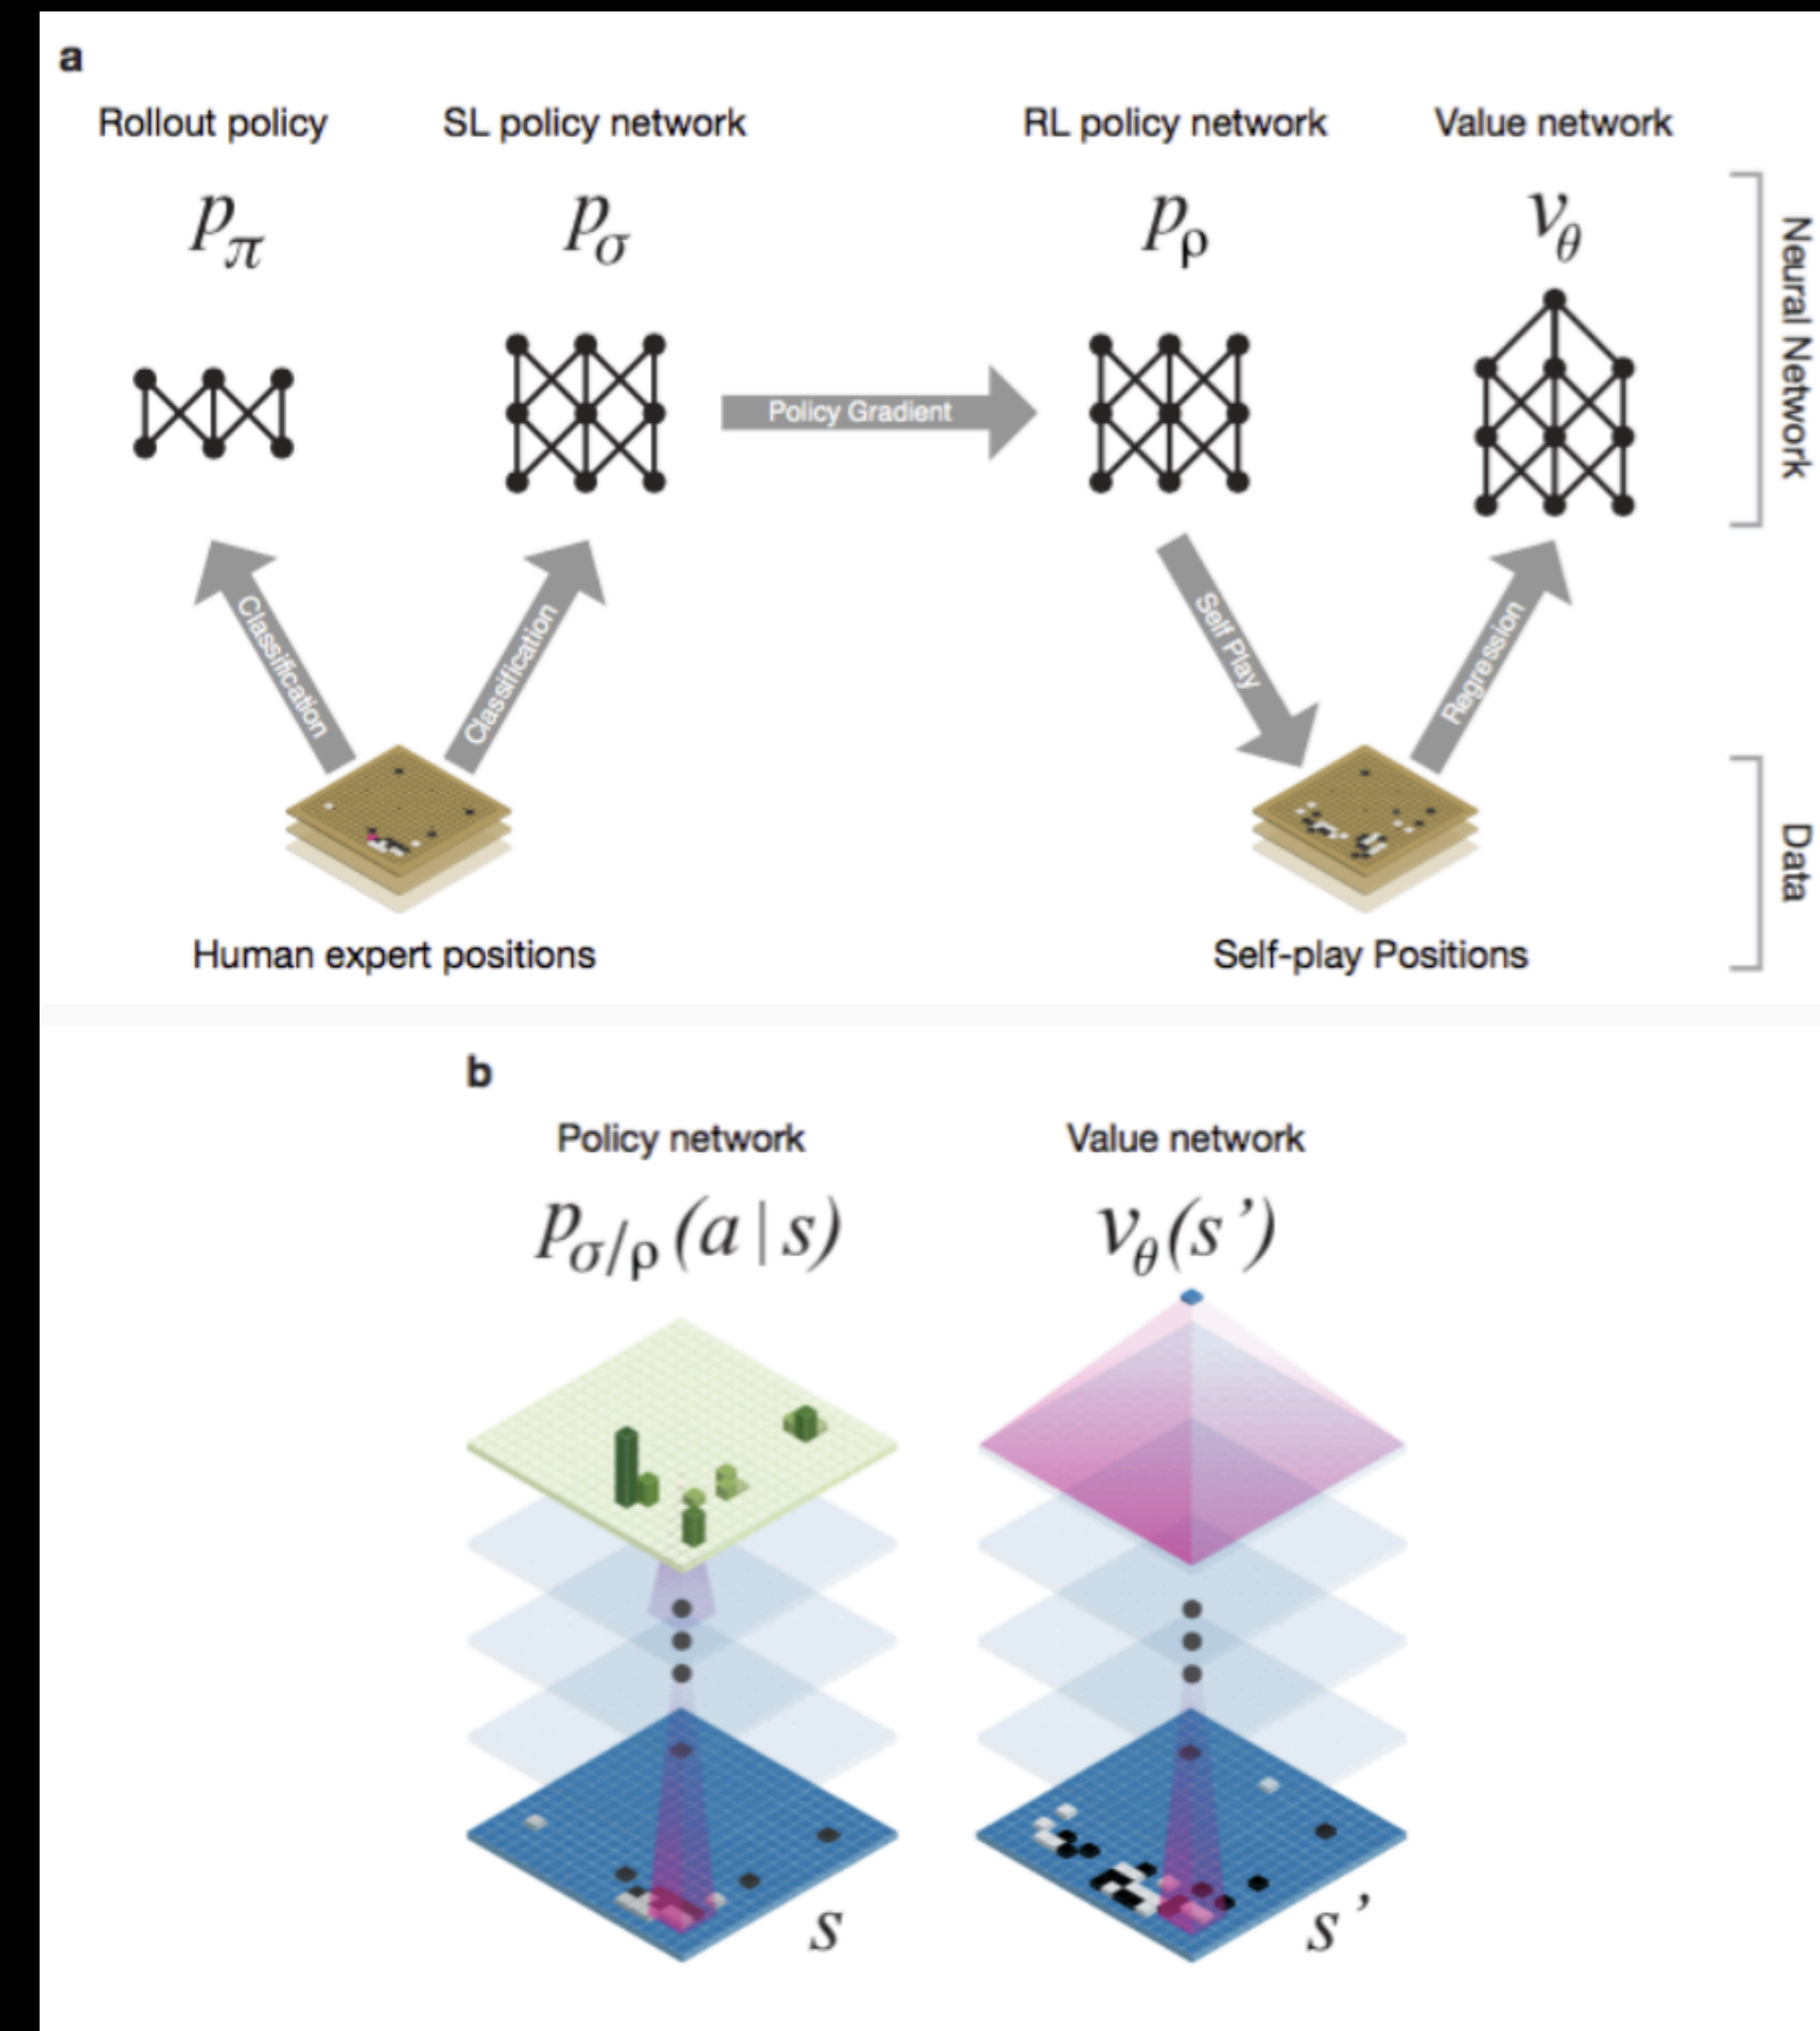
- ▶ Evaluation functions are always imperfect
- ▶ The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- ▶ An important example of the tradeoff between complexity of features and complexity of computation



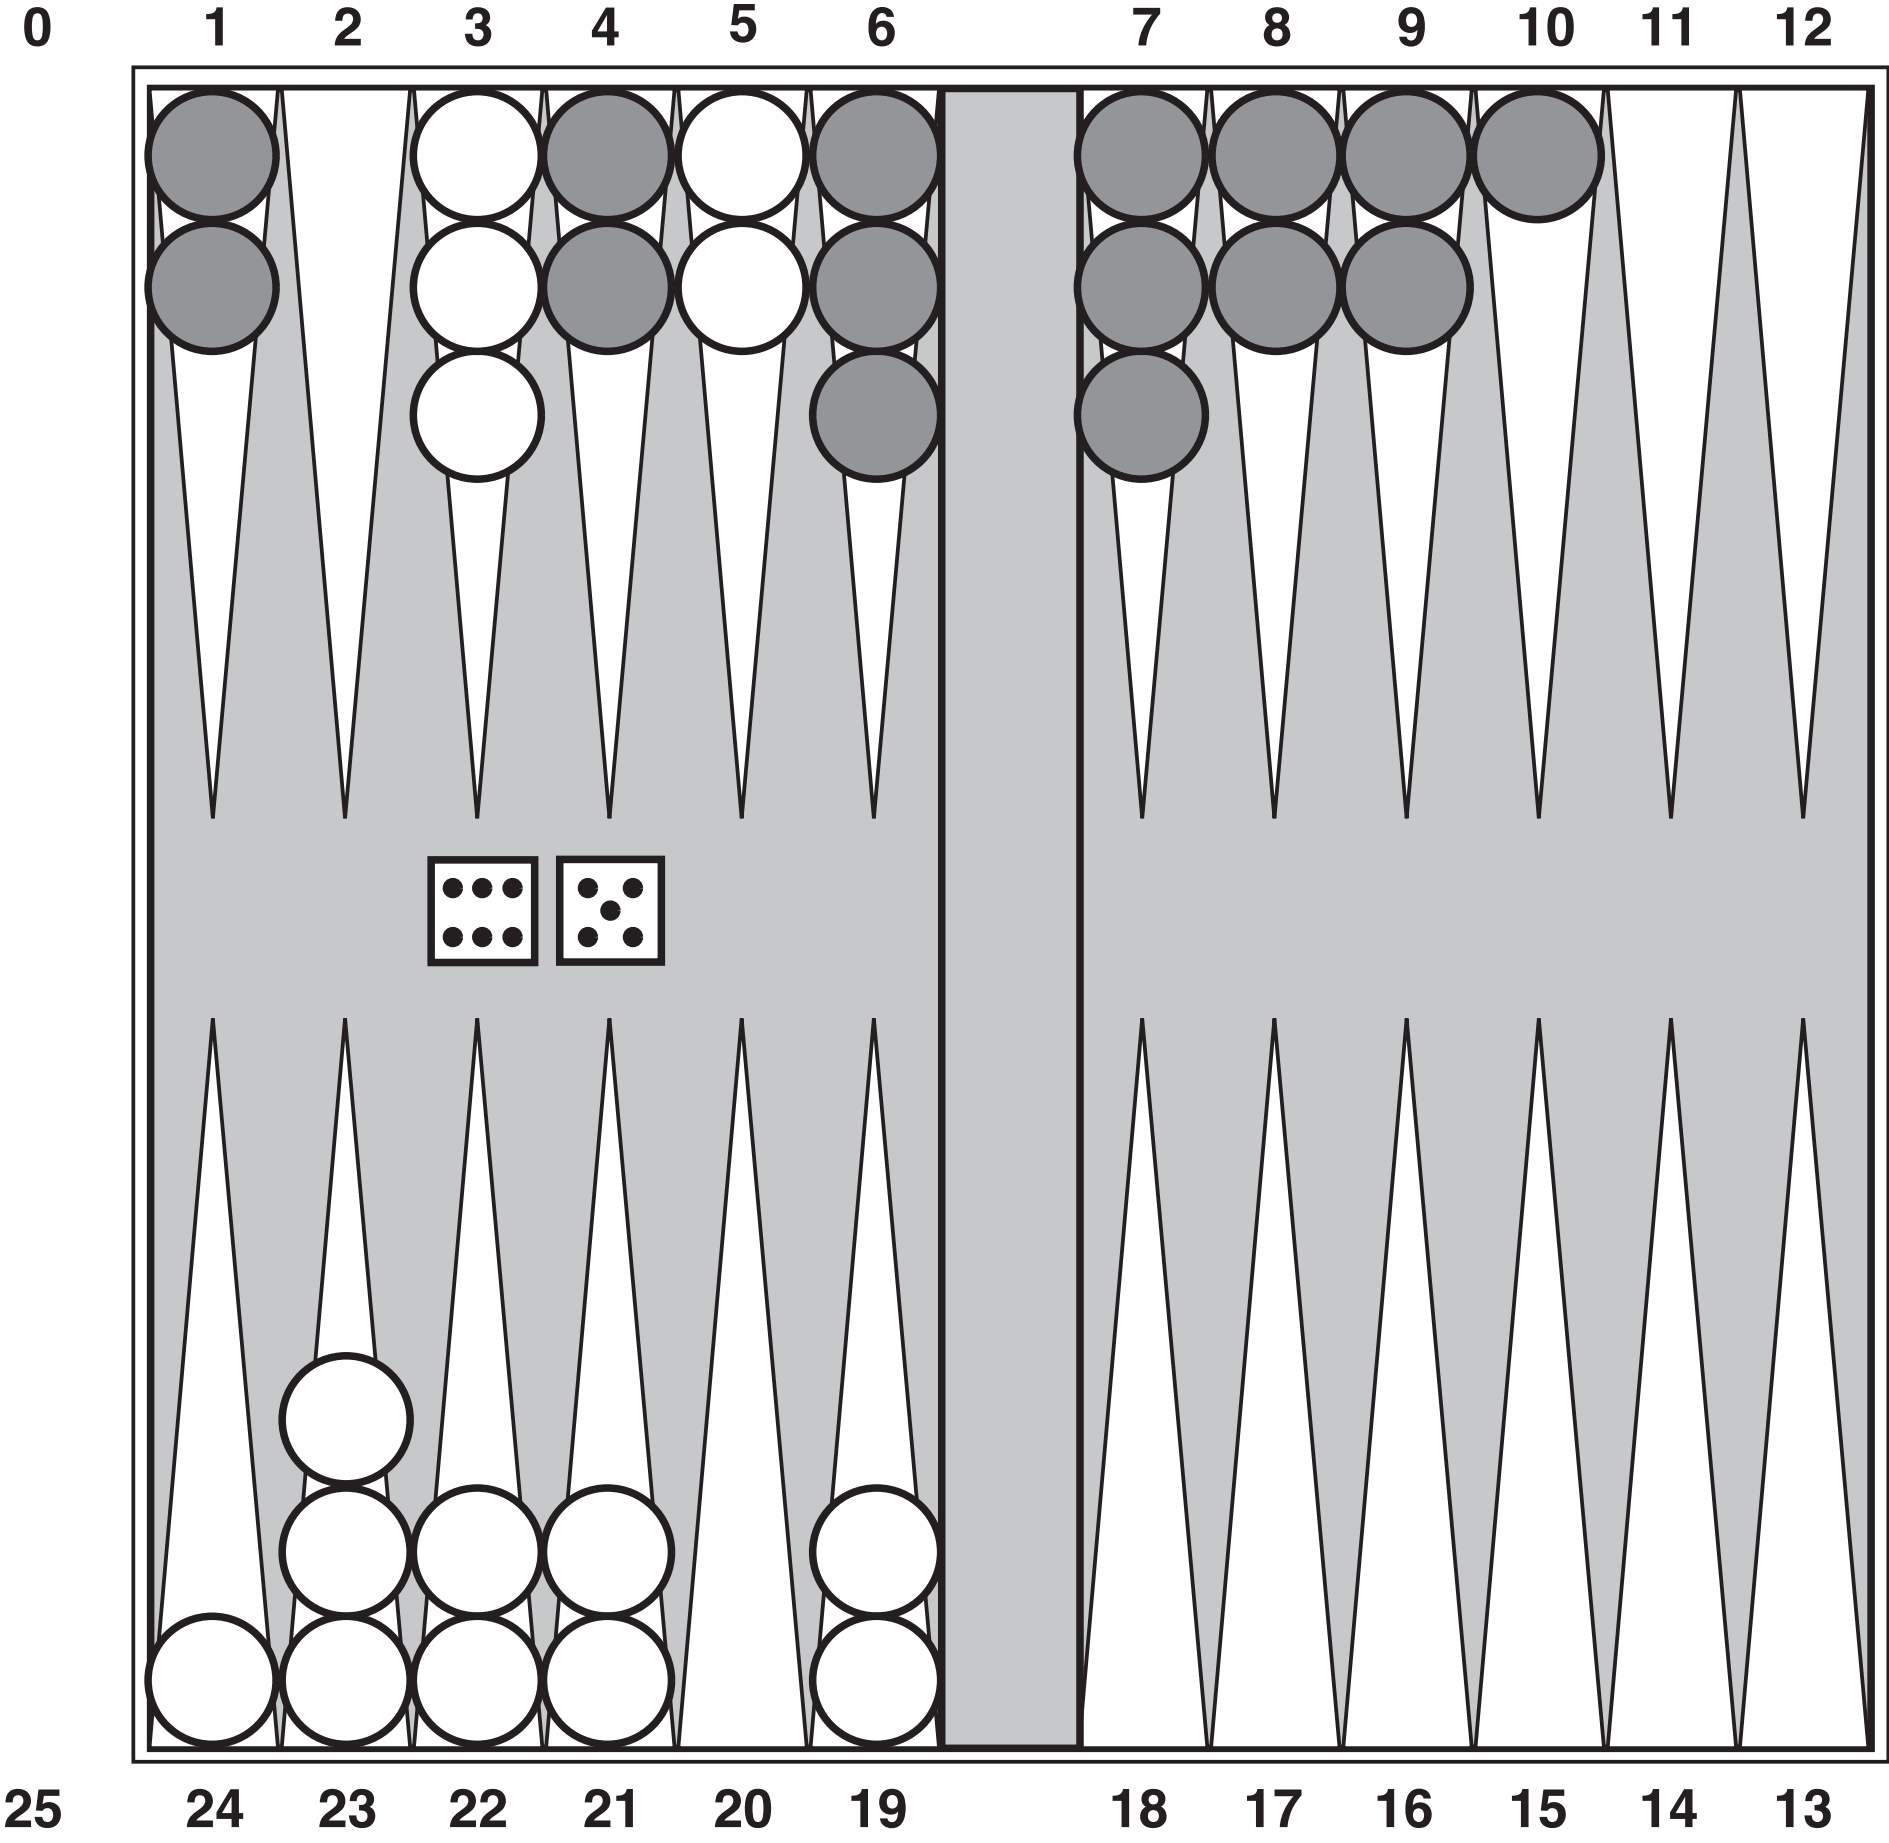
**HOW COULD YOU LEARN A
GOOD EVALUATION FUNCTION?**

ALPHAGO SYSTEM

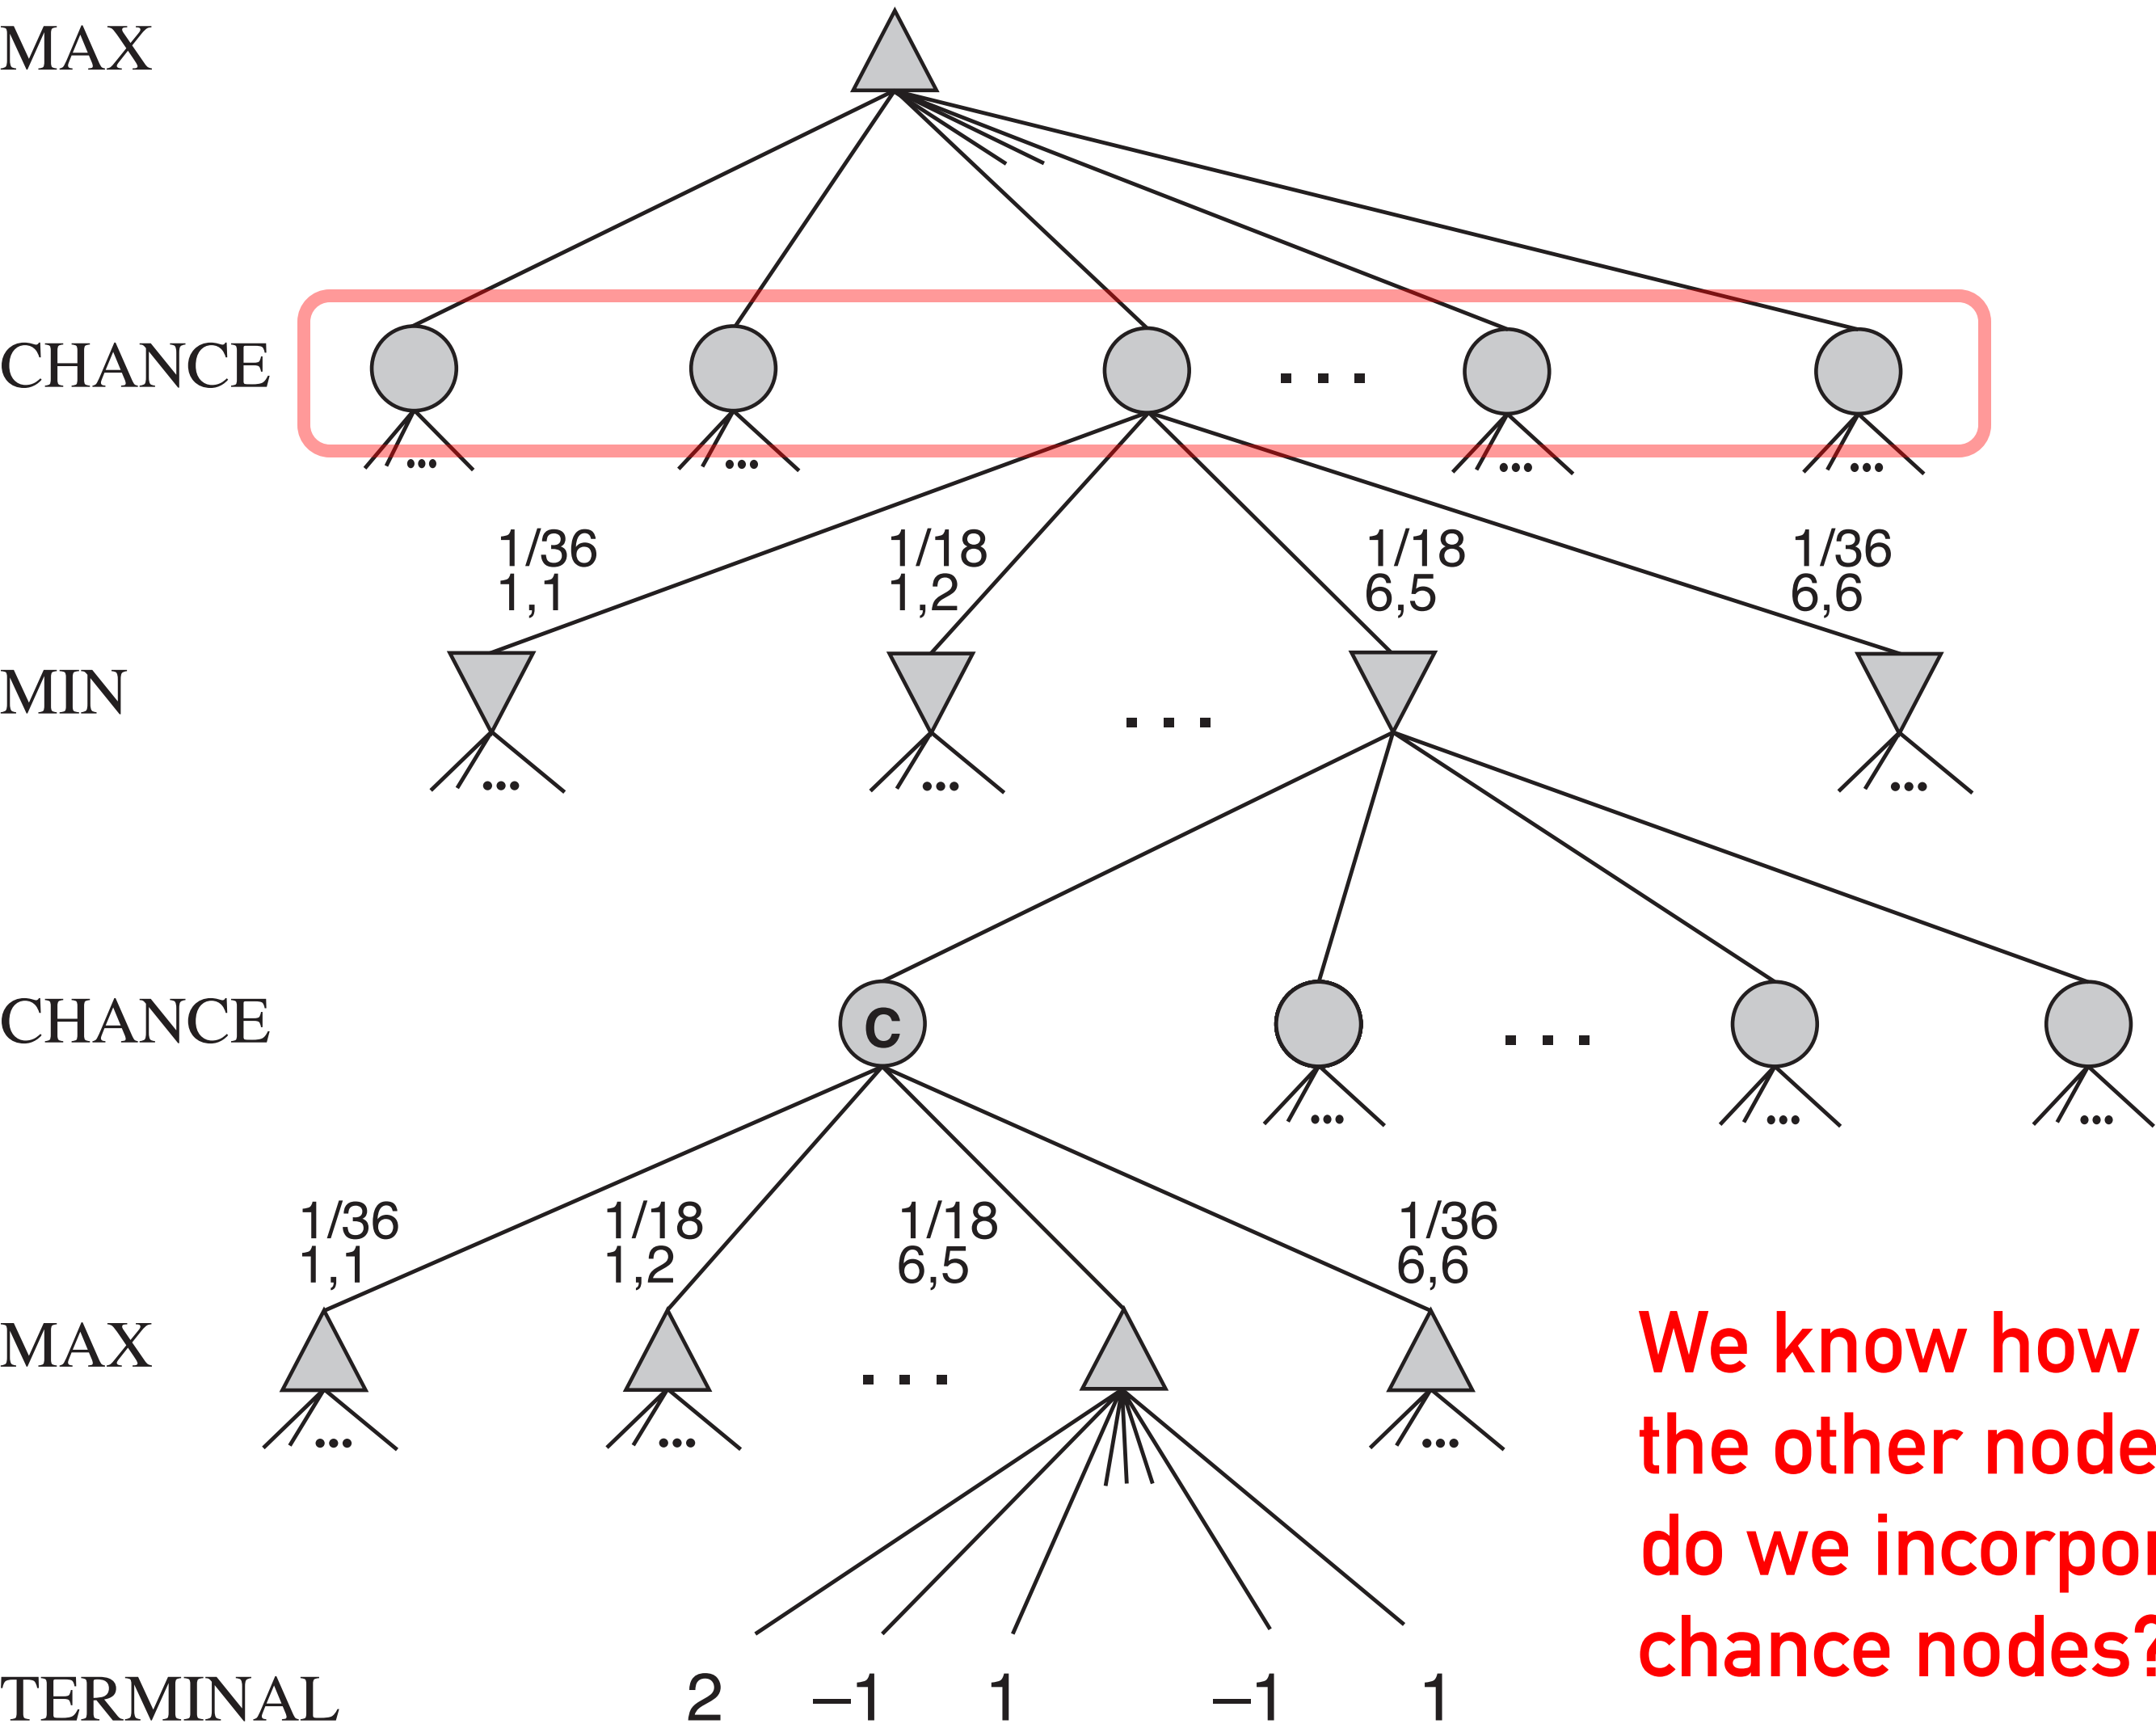
- **Deep learning + Reinforcement learning**
- One model predicts next move, given current state of board, *trained on 30 million positions from human games*
- Another model predicts likelihood of winning given current state, *trained on 30 million positions from self-play*
- **System** combines two models using Monte Carlo search



WHAT IF A GAME HAS A “CHANCE ELEMENT”?



GAME TREE WITH CHANCE ELEMENT



We know how to value the other nodes. How do we incorporate chance nodes?

EXPECTED VALUE

- ▶ The sum of the probability of each possible outcome multiplied by its value:

$$E(X) = \sum_i p_i x_i$$

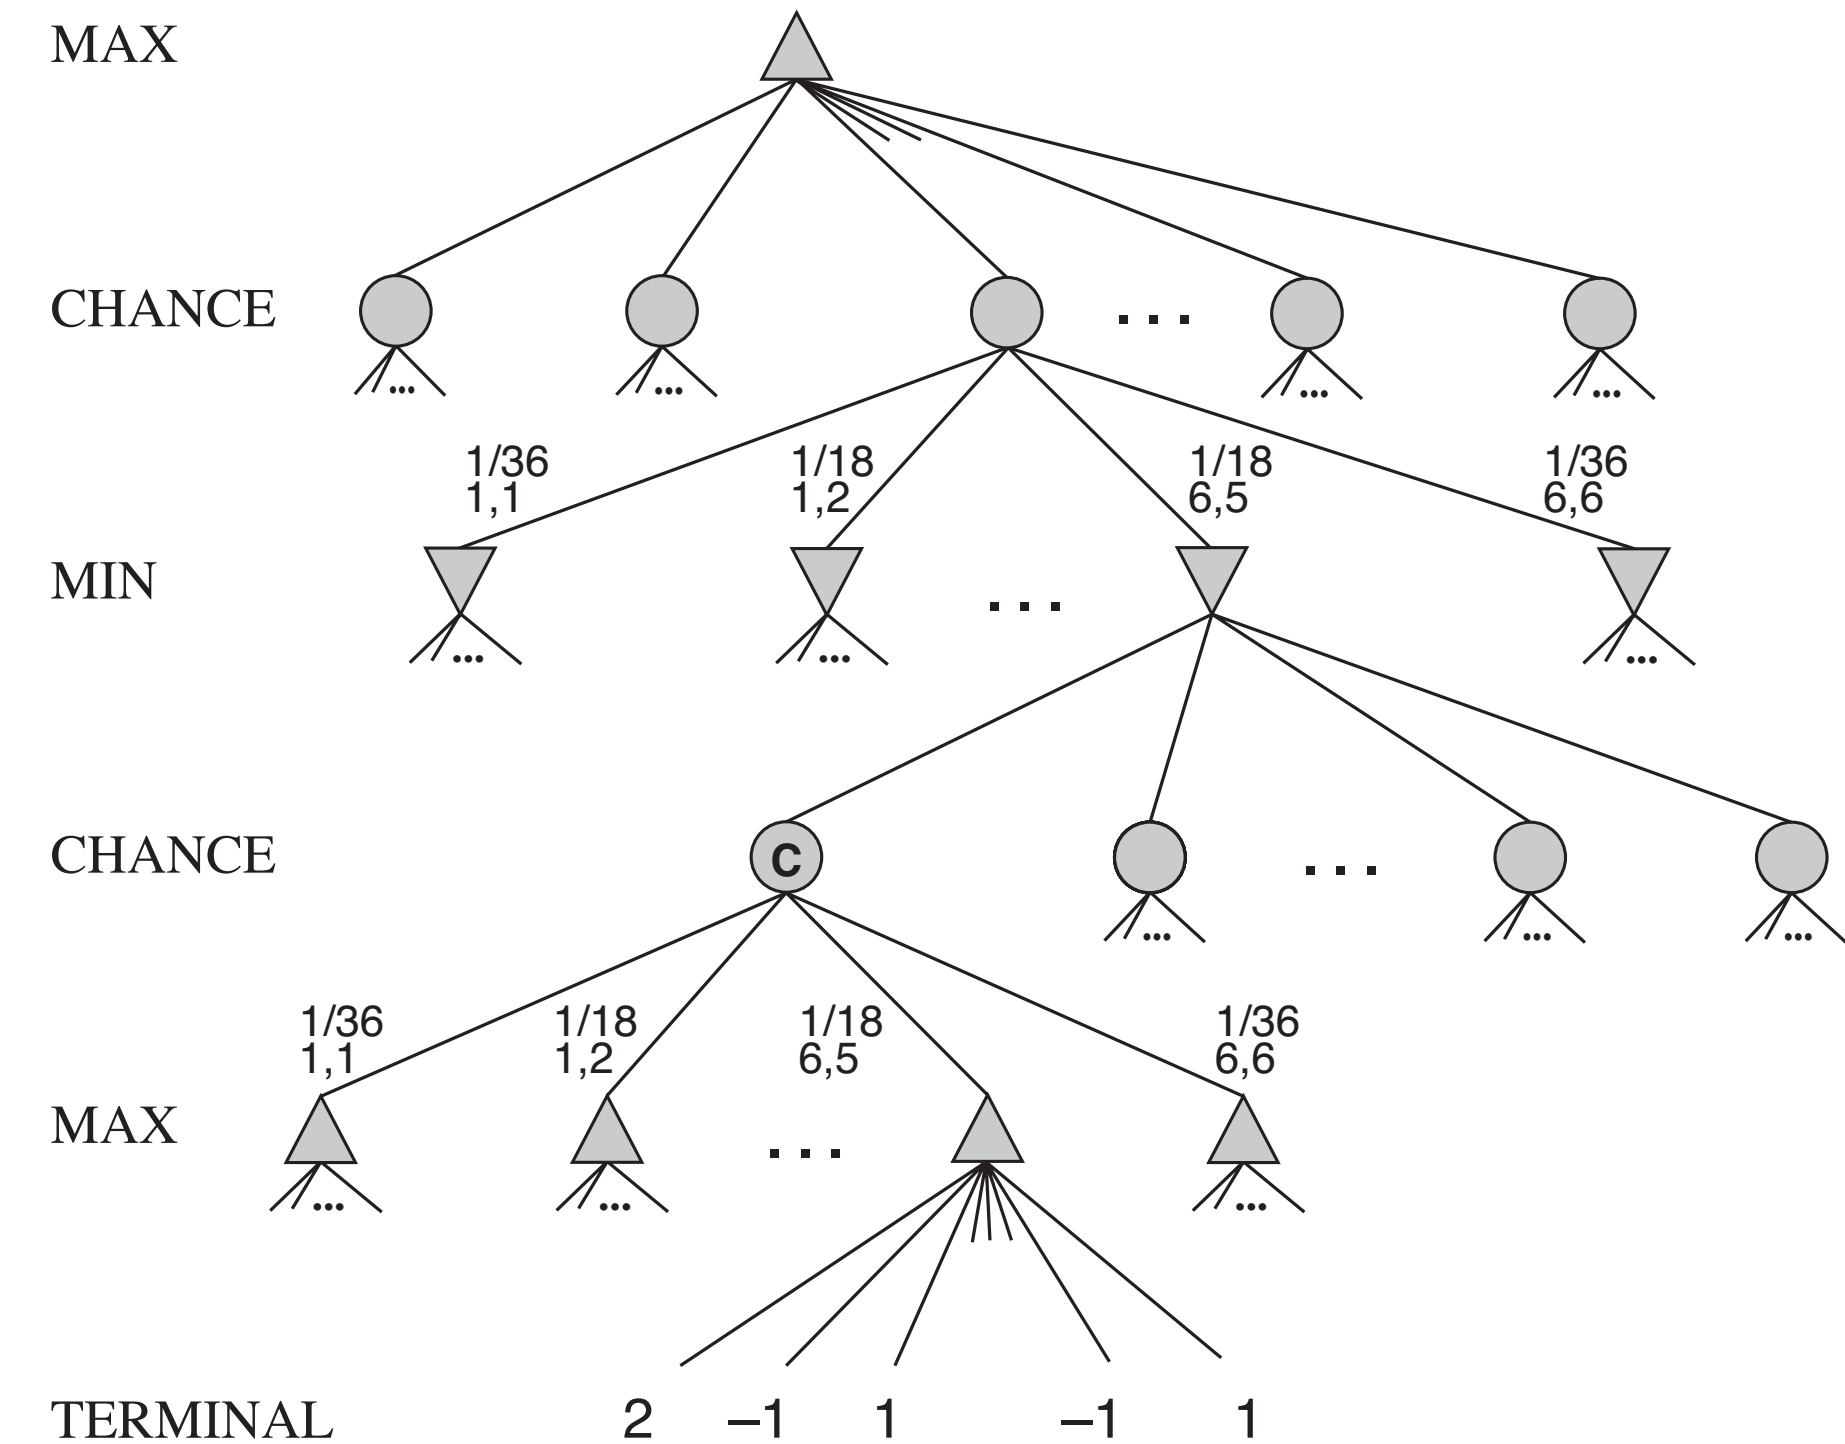
- ▶ Are there pathological cases where this statistic could do something strange
 - ▶ Extreme values ("outliers")
 - ▶ Functions that are a non-linear transformation of the probability of winning

EXPECTED MINIMAX VALUE

- ▶ Now three different cases to evaluate, rather than just two.


▶ MAX, MIN, CHANCE

- ▶ $\text{EXPECTED-MINIMAX-VALUE}(n) =$
 $\text{UTILITY}(n)$
 $\max_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$
 $\min_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$
 $\sum_{s \in \text{successors}(n)} P(s) \times \text{EXPECTED-MINIMAX}(s)$




if terminal node
 if MAX node
 if MIN node
 if CHANCE node

CONSTRAINT SATISFACTION PROBLEMS





WHAT IS SEARCH FOR?

- ▶ **Planning:** sequences of actions
 - ▶ The path to the goal is the important thing
 - ▶ Paths have various costs, depths
 - ▶ Heuristics give problem-specific guidance
- ▶ **Identification:** assignments to variables
 - ▶ The goal itself is important, not the path
 - ▶ All paths at the same depth (for some formulations)

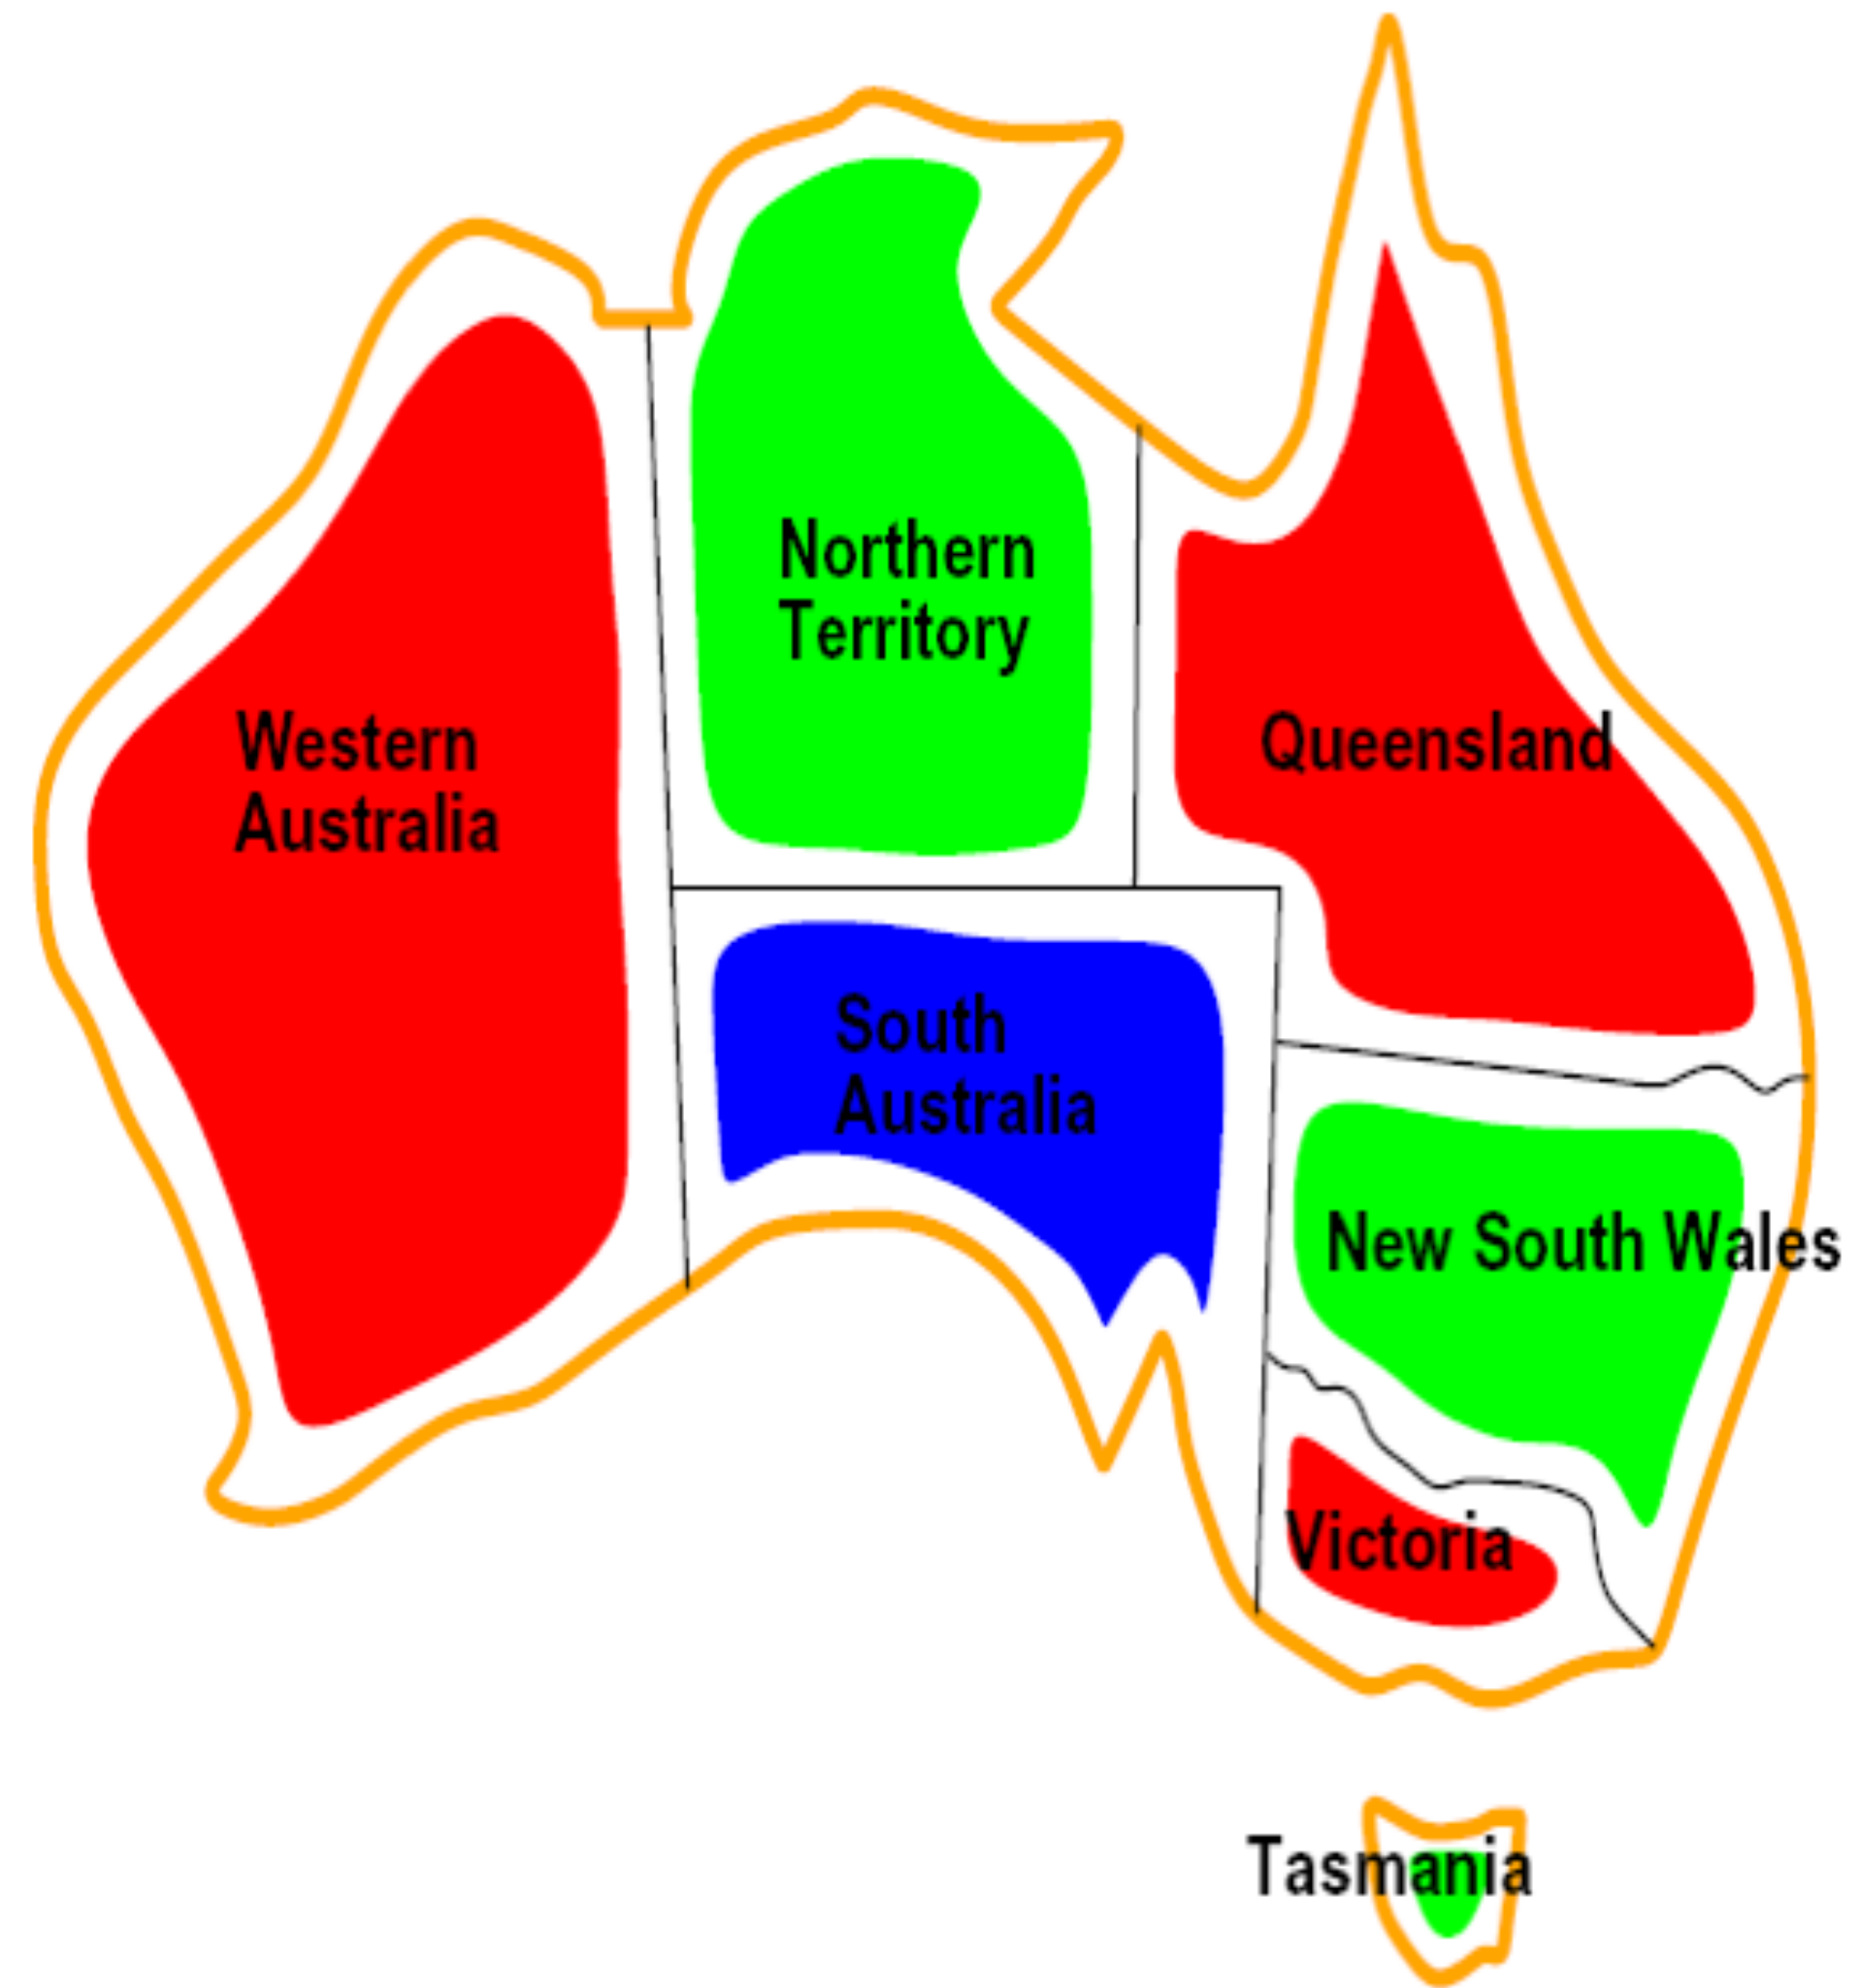


CONSTRAINT SATISFACTION PROBLEMS

- ▶ Standard search problems:
 - ▶ State is a “black box”: arbitrary data structure
 - ▶ Goal test can be any function over states
- ▶ Constraint satisfaction problems (CSPs) – a special subset of search problems
 - ▶ State is defined by **variables X_i** with values from a **domain D**
 - ▶ Goal test is a set of constraints specifying allowable combinations of values for subsets of variables
- ▶ Simple example of a formal representation language
- ▶ Allows useful general-purpose algorithms with more power than standard search algorithms



CSP EXAMPLES



EXAMPLE: MAP COLORING


- ▶ Variables: WA, NT, Q, NSW, V, SA, T
- ▶ Domains: $D = \{\text{red, green, blue}\}$
- ▶ Constraints: adjacent regions must have different colors

Implicit: $WA \neq NT$


Explicit: $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), \dots\}$

- ▶ Solutions are assignments satisfying all constraints, e.g.:

$\{WA=\text{red}, NT=\text{green}, Q=\text{red}, NSW=\text{green},$
 $V=\text{red}, SA=\text{blue}, T=\text{green}\}$




CONSTRAINT GRAPHS



CONSTRAINT GRAPHS

- ▶ Binary CSP: each constraint relates (at most) two variables
- ▶ Binary constraint graph: nodes are variables, arcs show constraints
- ▶ General-purpose CSP algorithms use the graph structure to speed up search.
E.g., Tasmania is an independent subproblem



EXAMPLE: N-QUEENS

► Variables: Q_k


► Domains: $\{1, 2, 3, \dots, N\}$

► Constraints:

Implicit: $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$

...



EXAMPLE: SUDOKU

- ▶ Objective
 - ▶ Fill the empty cells with numbers between 1 and 9
- ▶ Rules
 - ▶ Numbers can appear only once on each row
 - ▶ Numbers can appear only once on each column
 - ▶ Numbers can appear only once on each region
- ▶ Variables? Domain?
- ▶ Constraints?


8			4		6			7
						4		
	1					6	5	
5		9		3		7	8	
				7				
	4	8		2		1		3
	5	2					9	
		1						
3			9		2			5

SOLVING CSPS




STANDARD SEARCH FORMULATION

- ▶ Standard search formulation of CSPs
- ▶ States defined by the values assigned so far (ie. partial assignments)
 - ▶ Initial state: the empty assignment, $\{\}$
 - ▶ Successor function: assign a value to an unassigned variable
 - ▶ Goal test: the current assignment is complete and satisfies all constraints




SEARCH METHODS

- ▶ What would BFS do?
- ▶ What would DFS do?
- ▶ What problems does naive state space search have in this setting?



BACKTRACKING SEARCH



BACKTRACKING SEARCH

- ▶ Backtracking search is the basic uninformed algorithm for solving CSPs
- ▶ Idea 1: One variable at a time
 - ▶ Variable assignments are commutative, so fix ordering and only consider assignments to a single variable at each step
 - ▶ I.e., [WA = red then NT = green] same as [NT = green then WA = red]
- ▶ Idea 2: Check constraints as you go
 - ▶ “Incremental goal test” i.e. consider only values which do not conflict previous assignments
 - ▶ Might have to do some computation to check the constraints
- ▶ Depth-first search with these two improvements is called **backtracking search** (not the best name)
- ▶ Can solve n-queens for $n \approx 25$



BACKTRACKING EXAMPLE

